

# Documentation: `src/bot/graph.py`

---

This document explains in detail the **structure and functioning** of the `graph.py` module, which orchestrates an AI-driven conversation around **sleep-related health consultations**.

It integrates **LangGraph** state graphs with **LangChain** LLM prompt pipelines to enable smart, stateful, safe, and procedurally guided dialogues.

---

## 1. Purpose

The `graph.py` script defines a **state machine** for the chatbot that:

1. Starts conversations with a patient (optionally using a referral letter).
  2. Ensures patient input **stays on-topic** (sleep-related).
  3. Performs **safety checks** for self-harm risk.
  4. Dynamically decides whether to **ask questions** or **generate a professional summary**.
  5. Handles conversation flow persistently using a SQLite checkpointing system.
- 

## 2. Key Components

### 2.1 Dependencies

The script imports several key libraries and modules:

- **LangChain Core:**
    - `HumanMessage`
    - `AIMessage`
    - `ChatPromptTemplate`
  - **LangChain OpenAI:** `ChatOpenAI`
  - **LangGraph:**
    - `StateGraph` (for defining state transitions)
    - `END` terminal constant
    - `SqliteSaver` checkpoint backend
  - **Custom Project Modules:**
    - `GraphState` (state schema definition)
    - Pydantic data models: `GuardrailDecision`, `SuicideCheckDecision`, `SleepSummary`, `RouterDecision`
    - Prompt helper functions for system and human messages
  - **Persistence:** `sqlite3` to store conversation states
- 

## 3. LLM Configuration

Two LLM instances are initialized:

```
llm = ChatOpenAI(model="gpt-4o")
llm_summary = ChatOpenAI(model="gpt-4o", max_tokens=3000)
```

- **llm** → General dialogue & classification tasks.
  - **llm\_summary** → Summary generation with more tokens for detailed outputs.
- 

## 4. Graph Nodes

The LangGraph is based on **nodes** that mutate **GraphState**.

Nodes return **partial state updates**, merged into the persistent conversation state.

### 4.1 Guardrail Node

Ensures conversation stays **on-topic (sleep health)**.

- Uses structured LLM output (**GuardrailDecision**) for classification.
- Maintains an **off\_topic\_counter**.
- After 3 off-topic responses, terminates with a warning.

#### Prompt Components:

- System: **get\_guardrail\_prompt()**
  - Human: User message + classification instructions.
- 

### 4.2 Suicide Check Node

Performs a **safety check** to detect self-harm/suicidal intent.

- Looks at last 5 messages.
- Structured output: **SuicideCheckDecision**.
- If **risk detected** at medium/high/immediate → Terminates conversation.
- Custom safety messages based on severity.

#### Prompt Components:

- System: **get\_suicide\_check\_prompt()**
  - Human: Last 5 messages as conversation context.
- 

### 4.3 Ask Question Node

Asks **sequential consultation questions**.

- If starting conversation → Uses **get\_initial\_question\_prompt()** with referral letter context.
- Else → **get\_followup\_question\_prompt()** with conversation history.

#### Prompt Components:

- System: **get\_ask\_question\_system\_prompt()**

- Human: Initial or follow-up prompt.

Tracks:

- `last_question` for guardrail recovery.
  - Increment `questions_answered` when on-topic.
- 

## 4.4 Summary Node

Generates **summaries** for both the **patient** and **doctor**.

2 Modes:

1. **Initial summary** (after enough questions, confirmation pending).
2. **Final summary** (after patient modifications).

Structured output: `SleepSummary` with:

- `doctor_summary`
- `patient_summary`
- `urgency_level`

Fallback: If structured output fails → plain LLM text.

---

## 5. Router Logic

Controls flow between:

- `"ask_question"`
- `"generate_summary"`

**Logic Steps:**

1. If summary already confirmed → `generate_summary`.
  2. If < 5 questions answered → Continue questions.
  3. Else → Let `RouterDecision` AI decide (context-aware).
- 

## 6. Termination Check

`should_terminate` function:

- Ends conversation if `terminate_reason` exists in state.
- 

## 7. Graph Construction

```
graph_builder = StateGraph(GraphState)

graph_builder.add_node("guardrail", guardrail_node)
```

```
graph_builder.add_node("suicide_check", suicide_check_node)
graph_builder.add_node("ask_question", ask_question_node)
graph_builder.add_node("summary", summary_node)
graph_builder.add_node("router", router_node)
```

### Edges:

- entry → **guardrail**
- guardrail → suicide\_check OR END
- suicide\_check → router OR END
- router → ask\_question OR summary
- ask\_question → END (await new input)
- summary → END OR continue for edits

---

## 8. Persistence

Uses `SqliteSaver` to store full conversation state in `conversations.sqlite`, keyed by `thread_id`.

---

## 9. Conversation Loop

The CLI entry point (`main_loop`):

1. Ask for `user_id`.
2. Optionally request referral letter.
3. Start conversation with personalized or default greeting.
4. Resume past sessions if no new referral provided.
5. Stream graph events until conversation termination.

---

## 10. Prompts

Prompts are **template-driven**, combining:

- **System Messages:** Define AI's role & guidelines.
- **Human Messages:** Context (conversation history, referral letter).

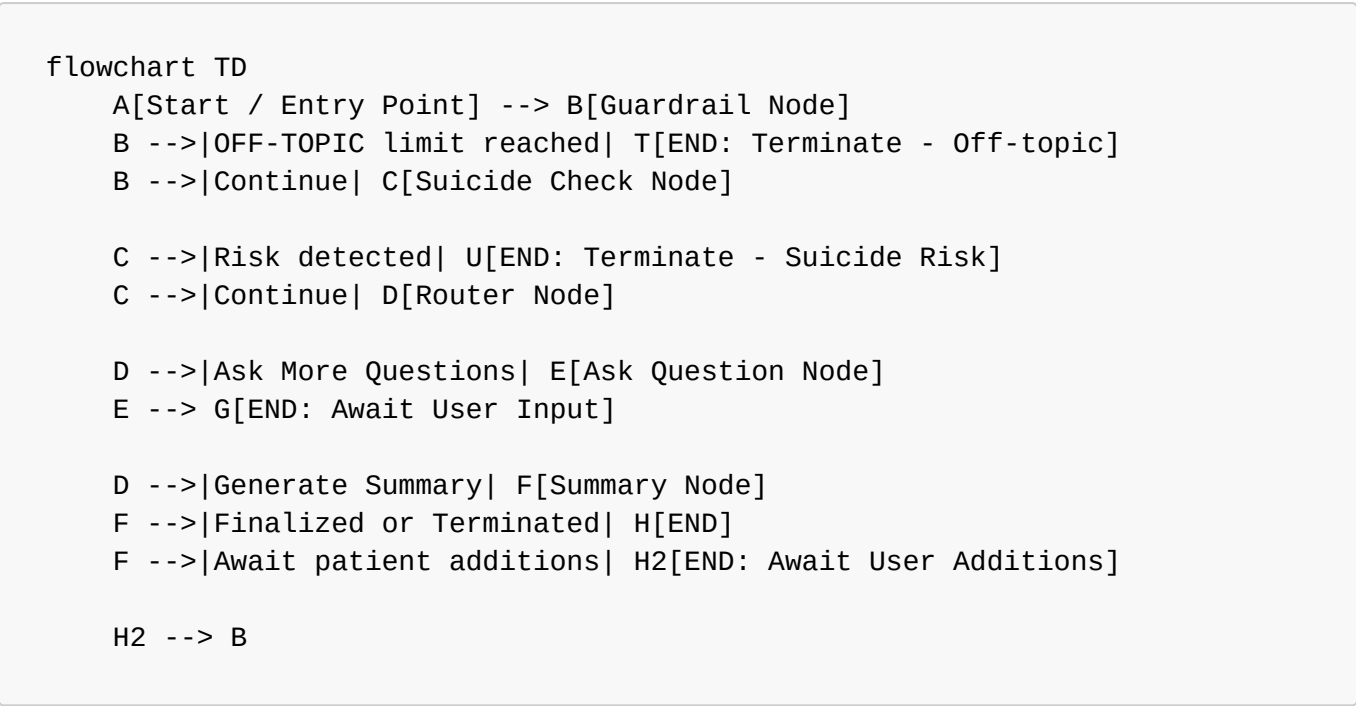
They are provided by `helper.py` functions for consistency.

---

## 11. Key Takeaways

- **Safety-first** → Guardrail & Suicide Check always run first.
  - **Dynamic flow** → Router logic + patient input determine direction.
  - **Persistent** → Conversations resume mid-way even after interruptions.
  - **Contextualized** → Referral letters greatly improve question relevance.
  - **Dual output** → Summaries tailored both for medical professionals and patients.
-

## 12. Visual Flow Diagram (Detailed)

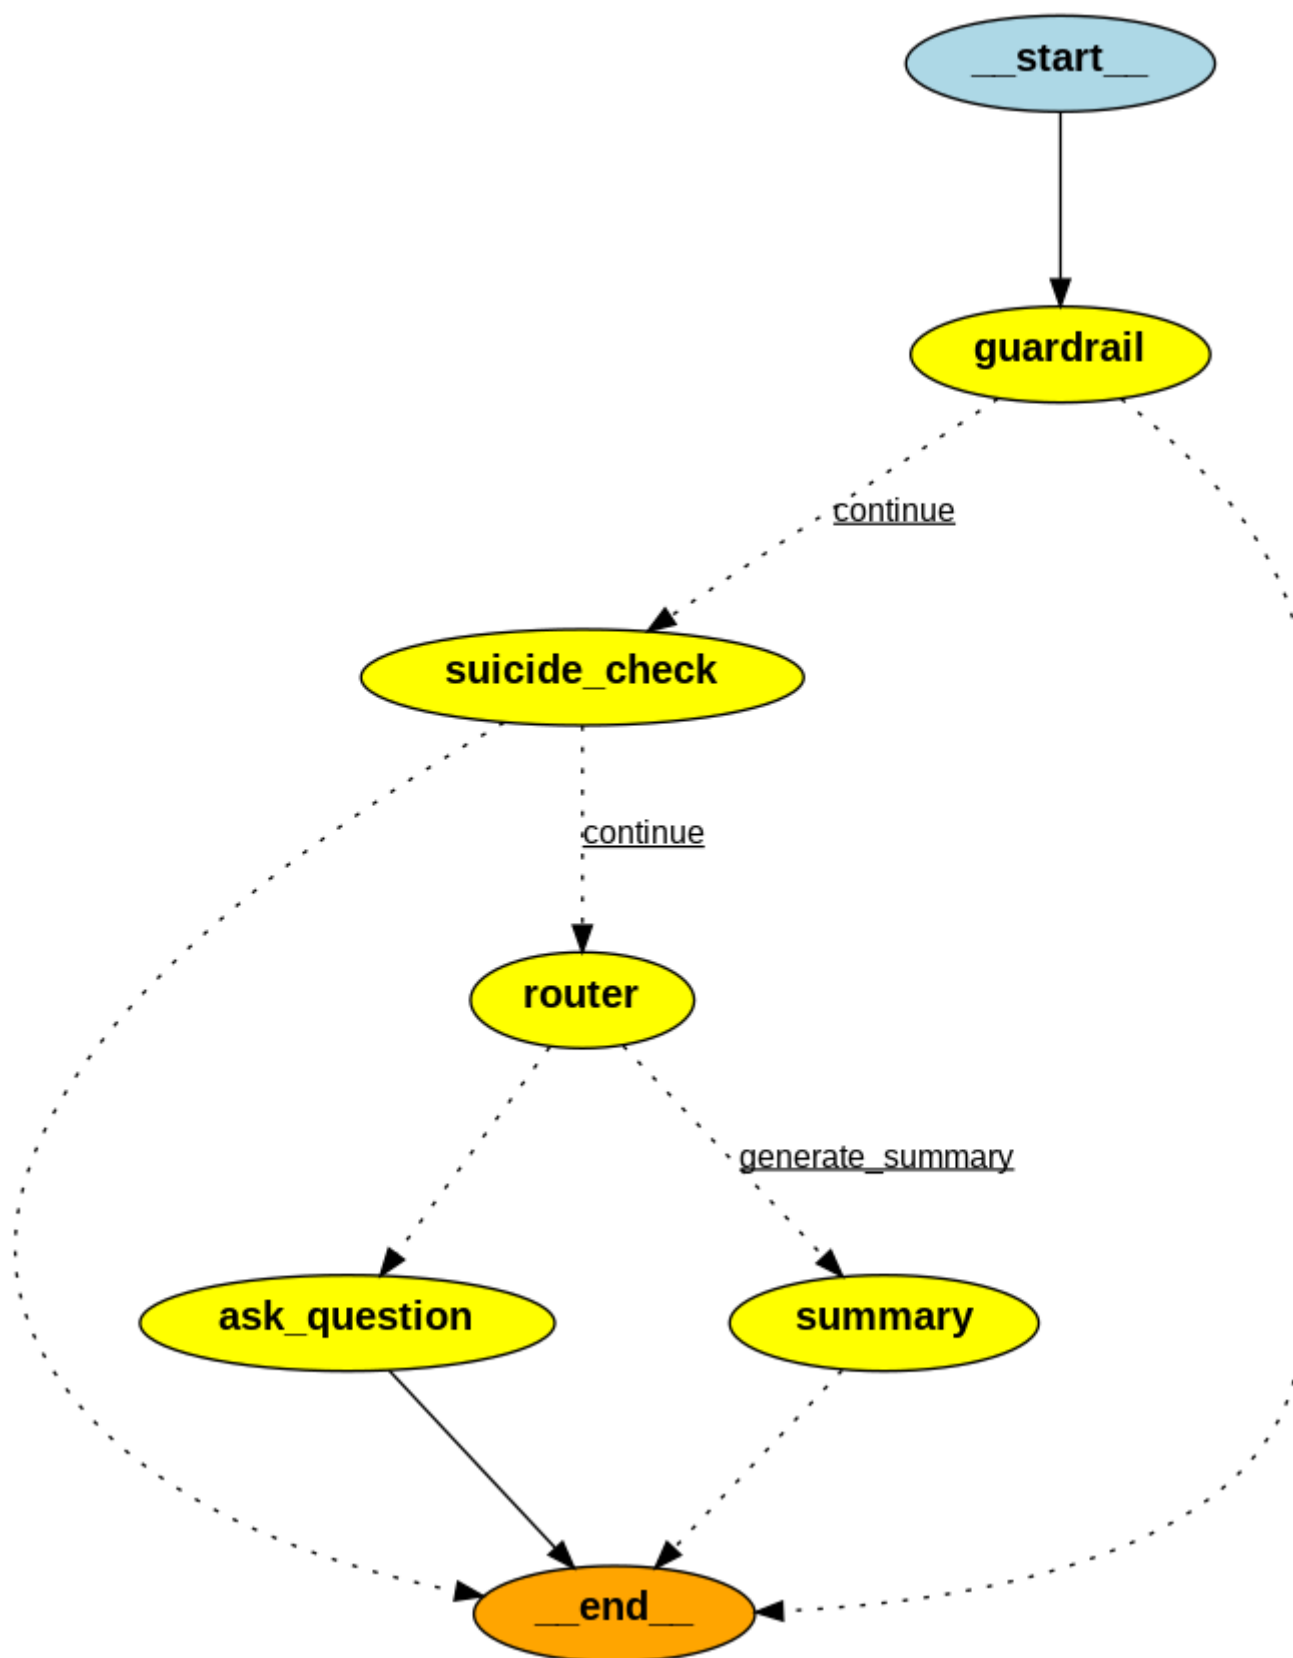


This **Mermaid diagram** illustrates the path branching and termination logic.

---

## 13. Graph Visualization (Generated)

Below is the actual **graph image** generated when running `graph.py` as the main module:



This PNG is built dynamically by `app.get_graph().draw_png()` during startup, ensuring it matches the code's current logic.

Maintained across nodes; can include:

- `messages` (chat history)
- `off_topic_counter`
- `questions_answered`

- last\_question
  - terminate\_reason
  - summary\_confirmed
  - referral\_letter
  - doctor\_summary
  - patient\_summary
  - urgency\_level
- 

## 14. Extensibility

Possible improvements:

- Add multilingual support in prompts.
- Save LLM reasoning metadata for audits.
- Include sentiment analysis tracking.